

# Painless Stochastic Gradient Interpolation, Line-Search, and Convergence Rates

NeurIPS 2019



Sharan  
Vaswani  
(Mila)



Aaron  
Mishkin  
(UBC)



Issam Laradji  
(UBC, EAI)



Mark  
Schmidt  
(UBC)



Gauthier  
Gidel  
(Mila, EAI)



Simon  
Lacoste-Julien  
(Mila)

# SGD is awesome, but...


Objective  $\min f(w) = \frac{1}{n} \sum_{i=1}^n f_i(w)$

Stochastic Gradient  
Descent (SGD)

$$w_{k+1} = w_k - \eta_k \nabla f_{ik}(w_k)$$

- + Simple to implement and use.
- + Good generalization properties.
- Slow convergence.
- Need to carefully tune the step-size.

**Painful ! :(**



# Making SGD painless

## Ingredient 1: Interpolation

- Satisfied by large *over-parametrized* models including typical neural networks and expressive kernel mappings.

$$||\nabla f(w^*)|| = 0 \implies \forall i, ||\nabla f_i(w^*)|| = 0$$

- Results in fast convergence of constant step-size SGD.

## Ingredient 2: Stochastic Line Search (SLS)

- Don't manually tune the step-size, automatically *search* for it.

Strong theoretical results and good empirical performance!

# SLS is simply SGD + Line search

In iteration  $k$ ,

1. Compute the gradients  $\nabla f_{ik}(w_k)$  for a given training batch
2. Search for a step-size  $\eta_k$  that satisfies the following *stochastic Armijo condition*,

$$f_{ik}(w_k - \eta_k \nabla f_{ik}(w_k)) \leq f_{ik}(w_k) - c \eta_k ||\nabla f_{ik}(w_k)||^2$$

3. Use the step-size and update the model parameters with SGD,

$$w_{k+1} = w_k - \eta_k \nabla f_{ik}(w_k)$$

# SLS is theoretically sound

Interpolation enables SLS



- Match *full-batch* convergence rates for smooth convex functions and strongly-convex functions.
- Achieve fast convergence rates for non-convex functions\*.

Other results:

- *Lipschitz line search* for Stochastic Extra-gradient (SEG) and its fast convergence for a subset of non-convex functions.
- Fast convergence of SEG + Lipschitz line search for a class of saddle point problems.

\* Additional assumptions apply

# SLS optimizes faster and better



ResNet-34 on CIFAR100

Other results: Deep matrix factorization, Kernels, Min-max games

# Try SLS in your project

## 1. Install our optimizer

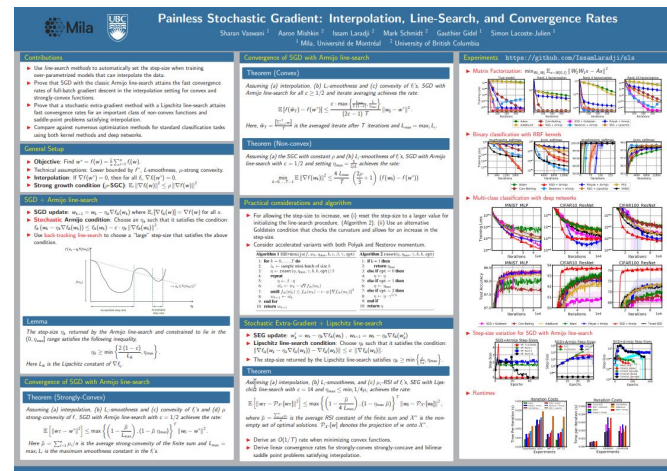
```
pip install --upgrade git+https://github.com/IssamLaradji/sls.git
```

## 2. Check out our Github code:

<https://github.com/IssamLaradji/sls>

## 3. Read the paper:

<https://arxiv.org/abs/1905.09997>



Come to our poster

Hall B + C #123

East Exhibition

Tue Dec 10th

5:30 PM